

University of Manchester

# Deriving Variables From Twitter/X Data: User Guide

Generated for replication as part of the DIGISURVOR project:  
<https://digisurv.or.github.io/main/>

**CITATION:** Gaughan, C., Gibson, R., Cantijoch, M., Cernat, A.,  
and Batista-Navarro, R., (2025). *Deriving Variables From Twitter/X  
Data: User Guide*. Version 1. Manchester: University of Manchester.

Dr. Conor Gaughan, Prof. Rachel Gibson, Dr. Marta  
Cantijoch, Prof. Alexandru Cernat, Dr. Riza Batista-  
Navarro  
9-19-2025

# INTRODUCTION

This user guide was created in conjunction with a working research paper: “*Linking Survey and Social Media Data: Bridging the Gap Between Open Research and Data Protection*”. It is the combined work of Dr. Conor Gaughan, Prof. Rachel Gibson, Dr. Marta Cantijoch, Prof. Alexandru Cernat, and Dr. Riza Batista-Navarro based at the University of Manchester, as part of the UKRI funded ‘smart data’ project DIGISURVOR. The guiding aim of the project is to investigate the methodological and ethical challenges of linking survey with social media data, and to develop ways in which researchers can better share their social media data with others. Data extracted from social media sites such as Facebook, Twitter (now X), Instagram, and LinkedIn are used frequently in social research, and it is becoming more common for these data to be linked with other data sources such as public opinion surveys. However, due to a range of legal, ethical and platform-specific restrictions, these data cannot be shared openly with other researchers. A particular challenge is that the publicly searchable nature of social media data means that releasing raw textual data risks disclosure of the author’s profile, which is especially problematic if linked to confidential survey data. Thus, in lieu of releasing raw social media data, we have established a procedure for the extraction of variables of interest from the data which can describe its structural and semantic content but remains unidentifiable.

The following user guide outlines this process which has been applied specifically to two linked survey-to-Twitter/X datasets of US respondents during the 2020 and 2024 US Presidential election campaigns. In this sense, the guide serves two purposes: 1) to act as a codebook for the two linked datasets which will be of use to researchers wishing to use this data for their own analysis; and 2) to act as a user guide for researchers who wish to replicate this process (or similar) on their own social media datasets. It first outlines the two datasets used as baselines for variable selection and creation including details about why and how they were

collected, their size and features, and some brief details about the linked survey component. The guide then describes the types of variables that have been extracted from this data, followed by descriptions of the required input variables and their data types from the raw data which are subsequently fed into the processing workflow. The guide then details the specific software requirements for running the code scripts which facilitate variable extraction, followed by details of the individual code scripts themselves. The final section provides individual descriptions of each of the 65 and 61 profile-level and tweet-level variables extracted from the raw input data.

## TWITTER (X) DATA:

Both datasets were collected for analysis of substantive research questions, as part of independent externally funded projects i.e. not the specific methodological questions posed in this project.<sup>1</sup> They were generated for a European Research Council funded project examining digital campaigning in five electoral democracies and combines two-wave pre- and post-election panel survey responses with participant Twitter/X data. Here, we draw on datasets for two Presidential election time points in the United States (US) – 2020 (1) and 2024 (2). The survey component measures media consumption, perceptions of digital campaign contact, awareness of misinformation, core political attitudes and behaviours, plus standard socio-demographic characteristics and self-reported Twitter/X use.

Wave 1 of Dataset 1 was fielded to 5,952 respondents between 16 September – 20 October 2020, and Wave 2 was fielded in the week following the election (9 November). Respondents to the pre-election survey were invited to share their Twitter/X handle with the research team. 2,460 participants reported having an active Twitter/X account (41%) and 1,598 agreed to share their handle (27% of overall sample; 65% of those with an account). There were four stages of attrition: 190 did not subsequently provide a handle, 171 provided an empty handle, and of the remaining 1,237, 920 could be validated against the Twitter/X API, which constituted 15% of the overall sample. Dataset 1 contains the collected timelines of these respondents during the period 1 September 2020 – 9 November 2020. Of the 920 respondents, we were able to extract profile-level variables for 905 and post-level variables for 539 of them respectively. Between the 539 respondents who had tweeted, there were 222,365 posts in total.

---

<sup>1</sup> Both were originally collected by Prof. Rachel Gibson for the European Research Council funded Project 833177(2020- 2025) Digital Campaigning and Electoral Democracy (DiCED).

Wave 1 of Dataset 2 was fielded to 5,757 respondents between 24<sup>th</sup> September and 21<sup>st</sup> October 2024 and Wave 2 was fielded in the two weeks following the election (12<sup>th</sup>-26<sup>th</sup> November 2024). Similar attrition rates were recorded in this survey: 2,621 had an active Twitter/X account (46%) and 1,570 agreed to share their handle (27% of overall sample; 60% of those with an account). There were four stages of attrition: 391 did not subsequently provide a handle and of the remaining 1,179, 963 could be validated against the Twitter/X API, which constituted 17% of the overall sample. As with Dataset 1, we collected the timelines of these respondents during the period 10<sup>th</sup> September – 4<sup>th</sup> November. We extracted profile-level variables for all 963 respondents, and post-level variables for 375. Between the 375 that had tweeted, there were 46,232 posts.

## VARIABLE CREATION:

We divide the SMD into two units of analysis: (1) Profile-Level; (2) Post-Level. We then divide the variables we want to derive from these units into two types: **Structural** and **Semantic**. We define structural variables as those which broadly *describe* the general content and structure of a respondent's digital data. This can include descriptive metrics such as a respondent's length of membership on the platform, their total number of followers, the frequency of their posting, their overall use of attachments and URLs in their posts, counting the number of characters in the text of a post, or the average amount of engagement a post receives. We define semantic variables as more computationally complex variables such as topic classification, lexical grammaticality and sophistication, language readability, toxicity and sentiment analysis, and ideological bias detection.

Variable Level	Overall
Profile-Level	65
Post-Level	61
<b>Overall</b>	126

To create these variables, we enlist several out-of-the-box computational tools which seek to combine optimal reliability with maximum simplicity and replicability. Thus, we do not train any models of our own or build any new variables that rely exclusively on the data we have. Moreover, the variables we create are not designed to be completely prescriptive or exhaustive. Instead, the aim here is to build a transparent and replicable *template* which can provide the skeleton code for others to implement and build upon using their own SMD. Additionally, we acknowledge that due to the authors' domain specific knowledge as well as the inherently political nature of the datasets we are using, some of these variables may be slanted towards a more political science audience. However, we believe that most of these baseline variables are applicable to many fields, but we encourage researchers to tailor these variables or generate

new ones according to their specific data and research objectives. The datasets we are using are also primarily in English given that we are using a US dataset and so certain variables may have to be adjusted for use on non-English datasets. All variables are built using the programming language *Python* (*v.3.11.7*).

### *Metadata Variables*

These are perhaps the most straightforward variables we can extract from any SMD. In most cases, this involves simple count measures, categorical variables or Booleans to describe the general structure of a unit of SMD. For (our) post-level data, this includes categorising the post type (original, retweet, reply, quote), the date it was created (timestamp is truncated to avoid too much granularity), engagement metrics (like count, retweet count, reply count, quote count), Booleans for whether it contains any media attachments or geotags, and the post reply settings. At the user profile level, we generate similar metadata metrics such as profile created at date, profile metrics (follower count, following count, like count, post count, listed count), as well as Booleans for whether additional profile fields are populated (display name, location, description). Where these fields are populated, we extract additional variables including matching listed locations to official geo-location data using the Nominatim API accessed via Python's *geopy* module as well as cross-referencing display names with real first and last names using Python's *names-dataset* module.

### *Text Variables*

Practically all SMD contains text-based fields. Text is the lifeblood of social networking, microblogging or forum platforms such as Facebook, Twitter/X, Reddit and LinkedIn. For media-based platforms such as YouTube, Instagram, TikTok and Snapchat, multimodal data like images, videos and audio are much more central, though text is still an important aspect of these SMD. Due to its greater complexity and to the fact we are using Twitter/X data, we do

not focus on multimodal data in this paper (though we do believe this is an important area to expand on in future work). Text is perhaps the most identifiable unit of information within our SMD which cannot be released in its raw format without significant risk of disclosure. Thus, we derive an array of anonymised structural and semantic variables from the text at both the post and user level.

First, we provide a set of variables which describe the lexical composition of the text. This involves simple count metrics for the total number of characters, number of textual characters and the number of capital letters, as well as the number of letters, words, unique words, sentences, syllables, monosyllables, and polysyllables. This can provide researchers with valuable information about the lexical and syntactic complexity of a text. Additionally, we derive a grammaticality score based on the Python's *LanguageTool* AI Grammer Checker, as well as a set of readability scores (Flesch Reading Ease, McAlpine EFLAW Score, Reading Grade) derived using the *textstat* module. We also calculate count measures for the number of special characters that appear in a text including the number of exclamation marks, question marks, emojis, @mentions, hashtags, and URLs. We believe these variables can be applied almost universally to any SMD text at both the post and user profile level, though one may wish to tailor it to capture specific nuances depending on the data, language, and platform.

### *URL Variables*

Many SMD text contain URLs which are an important aspect of content sharing and information dissemination, especially on microblogging sites like Twitter/X. Thus, not only do we count the number of URLs that appears in a body of text, we also match the domain level of the URL to the Media Bias/Fact Check (MBFC) media bias website.<sup>2</sup> MBFC is a comprehensive resource which contains media bias and credibility scores for thousands of

---

<sup>2</sup> The Media Bias/Fact Check website can be found here: <https://mediabiasfactcheck.com/>

major media sources, journalists and politicians and can be a useful tool for assessing the ideological slant and factual credibility of the information that users share. It is an imperfect method as it will not contain scores for all URLs if the domain does not match to anything in the MBFC dataset (typically because it is not a news source or it is too local). However, where users have shared information from relatively sized news sources, we can derive useful information about ideological balance and quality of content sharing, as well as content engagement (where users have engaged with posts that contain URLs).

### *TweetNLP*

Finally, the following semantic variables are generated using the Python module *TweetNLP*.<sup>3</sup> TweetNLP is a free-to-use module which supports a suite of natural language processing (NLP) tasks fine-tuned to SMD text. The NLP tasks supported include sentiment analysis, emotion recognition, offensive language identification, topic classification, irony detection, and named entity recognition. These are all implemented out-of-the-box in Python to derive semantic variables from the SMD text to give researchers a broad level understanding of the kinds of things users talk about and how they talk about it.

---

<sup>3</sup> More details can be found here: <https://github.com/cardiffnlp/tweetnlp>

## INPUT VARIABLES:

Input columns must be named as follows and in the correct data type. The code currently accepts .pkl files to ensure that data types are maintained.

### PROFILE DATA:

Variable Name	Type	Description
ID	String	The unique user ID
created_at	String	Date and time the profile was created
followers_count	Integer	Number of account followers
following_count	Integer	Number of accounts followed
tweet_count	Integer	Number of tweets posted
listed_count	Integer	Number of lists the account appears on
display_name	String	Profile display name
screen_name	String	Profile screen name (username/handle)
location	String	Profile location field
description	String	Profile description field (bio)

### TWEET DATA:

Variable Name	Type	Description
ID	String	The unique user ID
tweet_id	String	The unique tweet ID
referenced_tweet_type	String	Type of tweet
created_at_x	String	Date and time the author's tweet was created
source_x	String	Source of posting (iPhone, Browser, etc.)
language_x	String	Language of post text
retweet_count_x	Integer	Number of author retweets
reply_count_x	Integer	Number of author replies
like_count_x	Integer	Number of author likes
quote_count_x	Integer	Number of author quotes
media_keys	String	String of unique keys for media attachments within the text (GIFs, Videos, Images etc.)
geo_coordinates	String	String of geo-coordinates (Long, Lat)
reply_settings_x	String	Author reply settings
created_at_y	String	Date and time the original tweet was created
source_y	String	Original source of posting (iPhone, Browser, etc.)
language_y	String	Original language of post text
retweet_count_y	Integer	Number of original author retweets
reply_count_y	Integer	Number of original author replies
like_count_y	Integer	Number of original author likes
quote_count_y	Integer	Number of original author quotes
tweet_text	String	Text of the tweet
entities_urls	String	The entities returned by the API. This parsed to extract the expanded URL.

# SOFTWARE REQUIREMENTS

All output variables were extracted using the programming language Python (v.3.11.7). The following modules are required:

<b>Natural Language Processing (NLP)</b>		
re	Regular expression operations	<a href="https://docs.python.org/3/library/re.html">https://docs.python.org/3/library/re.html</a>
textstat	Calculate statistics from text	<a href="https://textstat.org/">https://textstat.org/</a>
emoji	Analyse emojis in text	<a href="https://pypi.org/project/emoji/">https://pypi.org/project/emoji/</a>
language_tool_python	Wrapper for <a href="#">LanguageTool</a> : used for grammar checking	<a href="https://pypi.org/project/language-tool-python/">https://pypi.org/project/language-tool-python/</a>
tweetnlp	Collection of NLP models for analysing tweet text	<a href="https://github.com/cardiffnlp/tweetnlp">https://github.com/cardiffnlp/tweetnlp</a>
rapidfuzz	Fuzzy string matching	<a href="https://rapidfuzz.github.io/RapidFuzz/">https://rapidfuzz.github.io/RapidFuzz/</a>
nameparser (HumanName)	Parsing human names into their individual components	<a href="https://pypi.org/project/nameparser/">https://pypi.org/project/nameparser/</a>
tldextract	Parsing URLs strings	<a href="https://pypi.org/project/tldextract/">https://pypi.org/project/tldextract/</a>
<b>Language, Name and Location Detection</b>		
langdetect	Detecting the language of text	<a href="https://pypi.org/project/langdetect/">https://pypi.org/project/langdetect/</a>
geopy (Nominatim)	Python client for geocoding web services	<a href="https://geopy.readthedocs.io/en/stable/">https://geopy.readthedocs.io/en/stable/</a>
names_dataset	First and last name matching	<a href="https://pypi.org/project/names-dataset/">https://pypi.org/project/names-dataset/</a>
<b>Data Wrangling</b>		
pandas	Fast and flexible data wrangling	<a href="https://pypi.org/project/pandas/">https://pypi.org/project/pandas/</a>
<b>System and Utility</b>		
time	Time access and conversions	<a href="https://docs.python.org/3/library/time.html">https://docs.python.org/3/library/time.html</a>
random	Generate pseudo-random numbers	<a href="https://docs.python.org/3/library/random.html">https://docs.python.org/3/library/random.html</a>
argparse	Parse command line arguments	<a href="https://docs.python.org/3/library/argparse.html">https://docs.python.org/3/library/argparse.html</a>
multiprocessing	Process-based parallelism	<a href="https://docs.python.org/3/library/multiprocessing.html">https://docs.python.org/3/library/multiprocessing.html</a>
tqdm	Process bar	<a href="https://pypi.org/project/tqdm/2.2.3/">https://pypi.org/project/tqdm/2.2.3/</a>

# PROCESS EXECUTION

There are eight code scripts in total that comprise the variable extraction process which can be executed via the command line using the following script:

*derive\_social\_media\_variables.py*

with the accompanying arguments:

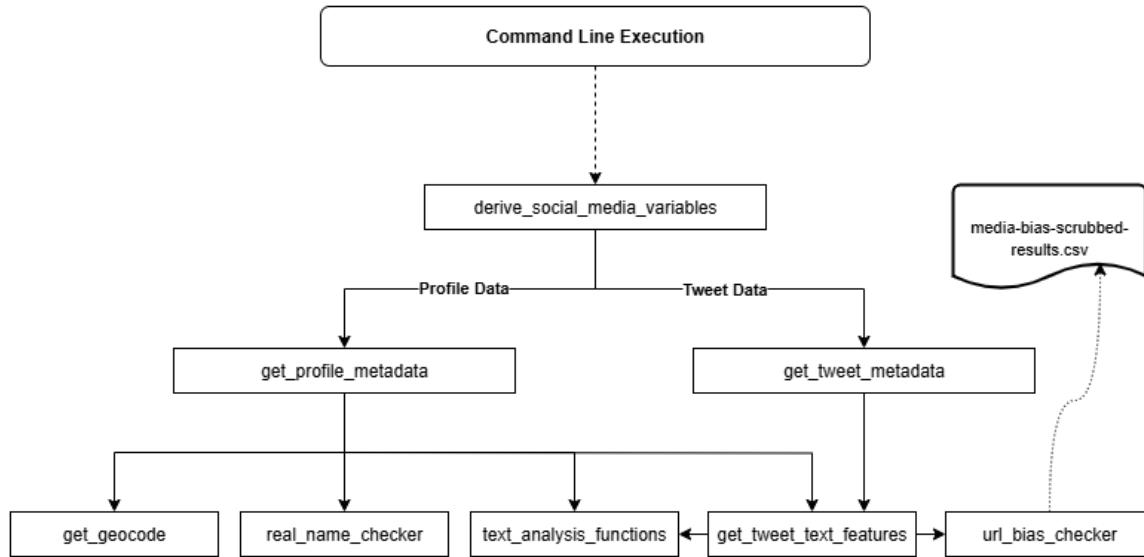
Argument	Description	Type	Required?	Options	Default
input	Path to input file	String	Yes	N/A	N/A
type	Input data type	String	Yes	‘profile’, ‘tweet’	N/A
file_path	Output file path (no extension!)	String	Yes	N/A	N/A
batch_size	Number of rows per batch	Integer	No	N/A	10
start_row	Row number to begin from	Integer	No	N/A	0
nproc	Number of processors to use	Integer	No	N/A	None
verbose	Enable verbose output	Store True	No	N/A	Defaults to True if argument is used

By default, the input file (profile\_data.pkl or tweet\_data.pkl), the data type (profile or tweet) and an output file path (automatically suffixed with “\_rows\_X\_to\_Y.csv”) are all required. By default, the input data provided is split into batches of 10 rows each for efficiency, but this can be optionally set higher or lower. Start row defaults to 0 but may be useful if the process breaks midway through after some rows have already been processed. The process itself can be relatively slow and process-based parallelism is relied upon to speed up the procedure, where batches of rows are processed in parallel to one another. The more processors that are used, the more rows that can be processed at once, although we recommend using less than the total number of processors available on the machine at once. If number of processors is not specified, it will default to the total number of processors on the machine minus one. If the verbose argument is specified, this will generate a process bar to track progress.

**Example:**

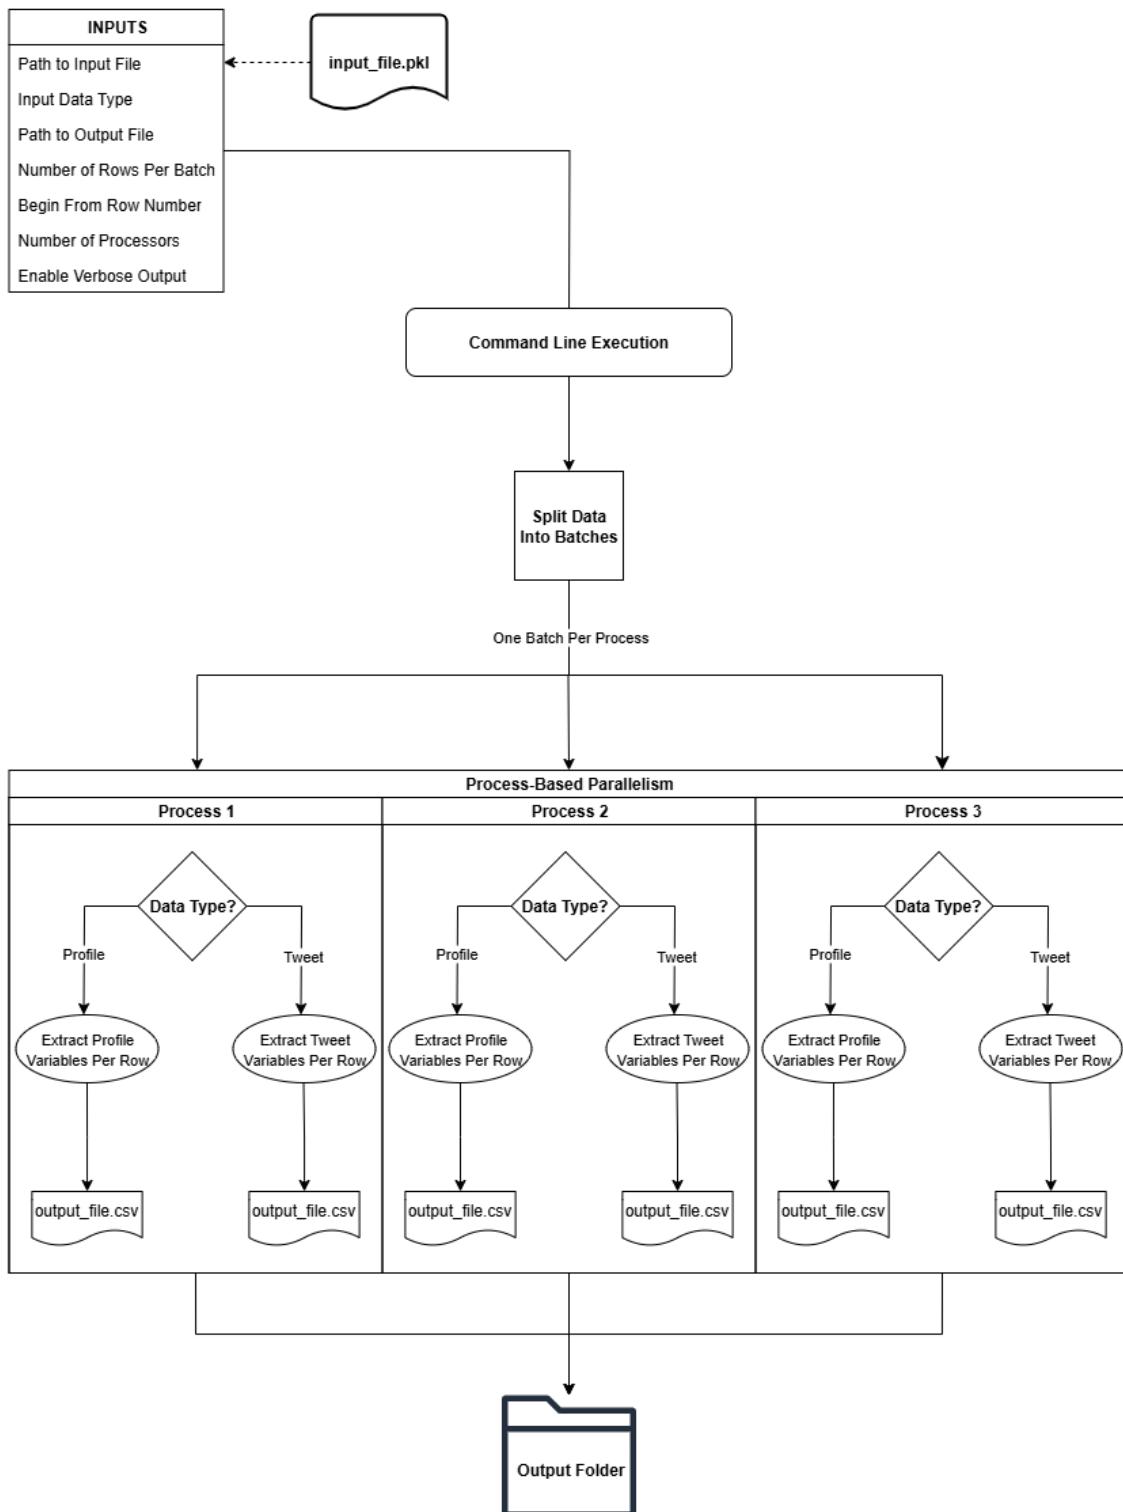
```
python.exe derive_social_media_variables.py --input input_file.pkl --type tweet --file_path  
tweet_variables --batch_size 1000 --start_row 0 --nproc 6 --verbose
```

# CODE STRUCTURE



The process relies on a total of eight unique code scripts. The `derive_social_media_variables` script is responsible for ingesting the input data file, splitting it into the required number of batches based on the specified batch size, and then executing each batch file in parallel via a designated processor. Where the number of overall batches exceeds the overall number of processors, these are queued until a previous batch is complete and a processor is freed. Depending on the data type, each processor will execute the `get_profile_metadata` or `get_tweet_metadata` scripts to extract new variables from the input data which is subsequently saved out as a .csv file to the designated file path. Ensure that all code scripts (and the accompanying media-bias .csv file) are all stored in the same directory (folder) along with an empty Python file titled `__init__.py`. This will mark the directory as a module and allow Python to import functions from the other scripts.

# DATA WORKFLOW:



# OUTPUT VARIABLES

## User-Level Variables

Variable Name	Description	Type	Software Dependency (Python Module)	Software Dependency (Module Function)
<i>ID</i>	A unique pseudonymised ID for the respondent	String		
<b>Profile Metrics:</b> These variables are returned directly by the platform API				
created_date	Date the account was created	String	datetime	date().strftime()
followers_count	Total number of followers	Integer		
following_count	Total number of accounts following	Integer		
tweet_count	Total number of tweets posted	Integer		
listed_count	Total number of lists the account appears on	Integer		
<b>Display Name (DN):</b> These variables are derived from the account display name field				
DN_length	Total number of characters	Integer	built-in	len()
The following variables were applied to pre-processed text. This involves the removal of URLs, emojis, hashtags, @mentions and extra whitespace. This is conducted using a custom-built strip_text() function which relies on the <i>re</i> and <i>emoji</i> modules. (Excl. the hashtag and emoji count variables)				
DN_character_length	Total number of textual characters	Integer	textstat	char_count()
DN_capt_character_count	Total number of capitalised characters	Integer	built-in	isupper()
DN_capt_character_prop	Proportion of capitalised characters to total number of characters	Float		
DN_letter_count	Total number of letters	Integer	textstat	letter_count()
DN_word_count	Total number of words	Integer	textstat	lexicon_count()

<i>DN_unique_word_count</i>	Total number of unique words	Integer	built-in	set()
<i>DN_exclamation_count</i>	Total number of exclamation marks	Integer	built-in	count()
<i>DN_question_mark_count</i>	Total number of question marks	Integer	built-in	count()
<i>DN_hashtag_count</i>	Total number of hashtags	Integer	re	findall()
<i>DN_emoji_count</i>	Total number of emojis	Integer	emoji	emoji_count()

For human name handling, a selection of functions were used. The `HumanName()` function from the `nameparser` module was used to parse listed names and salutations and first and last names were matched against the `NameDataset` from the `names_dataset` module: <https://pypi.org/project/names-dataset/>. The string similarity between a user's display name and their handle (username) is the Levenshtein Distance similarity ratio computed using the `ratio()` function from the `rapidfuzz` module.

<i>DN_handle_similarity</i>	Levenshtein Distance similarity between display name and handle (username)	Float	from rapidfuzz import fuzz	ratio()
<i>DN_listed_salutation</i>	Did the user list a salutation?	Boolean	from nameparser import HumanName	<code>HumanName().title</code>
<i>DN_used_first_name</i>	Did the user list a real first name?	Boolean	from names_dataset import NameDataset	search()
<i>DN_used_last_name</i>	Did the user list a real last name?	Boolean	from names_dataset import NameDataset	search()

**Location (LC):** These variables are derived from the account location field

<i>LC_field_populated</i>	Was the location field populated?	Boolean		
<i>LC_level_count</i>	Total number of location levels listed. This is ascertained by number of words separated by a comma or forward slash.	Integer	re	split(r'[,/]')

Geo-location matching is conducted through the Nominatim API which uses OpenStreetMap data for geocoding. It accepts free-form search queries and is accessed via the `geopy.geocoders` module.

<i>LC_levels_identified</i>	Total number of levels that can be matched to an official geographical location	Integer	from geopy.geocoders import nominatim	geocode()
<i>LC_level_types</i>	Location types identified	List of Strings	from geopy.geocoders import nominatim	geocode().raw.get("address_type")

**Description (DS):** These variables are derived from the account description field

<i>DS_field_populated</i>	Was the description field populated?	Boolean		
<i>DS_language</i>	Predicted text language	String	lang_detect	detect()
<i>DS_post_length</i>	Total number of characters	Integer	built-in	len()

The following variables were applied to pre-processed text. This involves the removal of URLs, emojis, hashtags, @mentions and extra whitespace. This is conducted using a custom-built strip\_text() function which relies on the *re* and *emoji* modules. (Excl. the hashtag, mention and emoji count variables)

<i>DS_character_length</i>	Total number of textual characters	Integer	textstat	char_count()
<i>DS_capt_character_count</i>	Total number of capitalised characters	Integer	built-in	is.upper()
<i>DS_capt_character_prop</i>	Proportion of capitalised characters to total number of characters	Float		
<i>DS_letter_count</i>	Total number of letters	Integer	textstat	letter_count()
<i>DS_word_count</i>	Total number of words	Integer	textstat	lexicon_count()
<i>DS_unique_word_count</i>	Total number of unique words	Integer	built-in	set()
<i>DS_sentence_count</i>	Total number of sentences	Integer	textstat	sentence_count()
<i>DS_syllable_count</i>	Total number of syllables	Integer	textstat	syllable_count()
<i>DS_monosyllable_count</i>	Total number of monosyllables	Integer	textstat	monosyllabcount()
<i>DS_polysyllable_count</i>	Total number of polysyllables	Integer	textstat	polysyllabcount()
<i>DS_grammar_score</i>	Proportion of grammatical errors to total number of words	Float	language_tool_python	tool_check()
<i>DS_reading_ease</i>	Flesch Reading Ease Score. The higher the score, the easier the readability of the text. (Max 121.22)	Float	textstat	flesch_reading_ease()

<i>DS_non_eng_reading_ease</i>	McAlpine EFLAW Readability Score. Readability of an English text for a foreign learner of English. It is recommended to aim for a score equal to or lower than 25	Float	textstat	mcalpine_eflaw()
<i>DS_reading_grade</i>	Estimated school grade level required to understand the text. Consensus score based on seven separate readability tests	Float	textstat	text_standard()
<i>DS_exclamation_count</i>	Total number of exclamation marks	Integer	built-in	count()
<i>DS_question_mark_count</i>	Total number of question marks	Integer	built-in	count()
<i>DS_hashtag_count</i>	Total number of hashtags	Integer	re	findall()
<i>DS_emoji_count</i>	Total number of emojis	Integer	emoji	emoji_count()
<i>DS_mentions_count</i>	Total number of @mentions	Integer	re	findall()
<i>DS_url_count</i>	Total number of URLs	Integer	re	findall()

Twitter/X reformats in-text URLs to its own domain: <https://t.co/>. To find the original URL domain, the full URL listed in the extended\_urls column must be used. The domain of URL is extracted using the extract().domain function from the *tldextract* module and matched to a list of domains in a pre-built CSV file scrubbed from the Media Bias/Fact Check website: <https://mediabiasfactcheck.com/>. This contains 1579 unique domains for major media sources around the world, as well as scores of their ideological bias and reporting credibility.

<i>DS_url_mbfc_matches</i>	Total number of URL domain matches against the Media Bias/Fact Check (MB/FC) database	Integer	built-in	Isin()
<i>DS_url_mbfc_bias</i>	Media Bias/Fact Check bias score(s) for URL domain matches	Integer		
<i>DS_url_mbfc_credibility</i>	Media Bias/Fact Check factual reporting credibility score(s) for URL domain matches	Integer		

The following variables are estimated using pre-trained models from the *TweetNLP* module: <https://github.com/cardiffnlp/tweetnlp>

<i>DS_topics</i>	Topic classification label(s)	String	<code>tweetnlp.load_model('topic_classification')</code>	<code>topic()</code>
<i>DS_topic_prob</i>	Certainty estimates for each topic label	Float		

<i>DS_sentiment</i>	Sentiment classification label	String	tweetnlp.load_model('sentiment')	sentiment()
<i>DS_sentiment_prob</i>	Certainty estimate for the sentiment classification	Float		
<i>DS_irony</i>	Irony classification label	String	tweetnlp.load_model('irony')	irony()
<i>DS_irony_prob</i>	Certainty estimate for the irony classification	Float		
<i>DS_offensive</i>	Offensive language classification label	String	tweetnlp.load_model('offensive')	offensive()
<i>DS_offensive_prob</i>	Certainty estimate for the offensive language classification	Float		
<i>DS_emotion</i>	Emotion classification label	String	tweetnlp.load_model('emotion')	emotion()
<i>DS_emotion_prob</i>	Certainty estimate for the emotion classification	Float		
<i>DS_hate</i>	Hate speech classification label	String	tweetnlp.load_model('hate')	hate()
<i>DS_hate_prob</i>	Certainty estimate for the hate speech classification	Float		
<i>DS_entity_count</i>	Total number of named entities recognised	Integer	tweetnlp.load_model('ner')	ner()
<i>DS_entity_types</i>	Types of named entities recognised	String		
<i>DS_entity_prob</i>	Certainty estimates for each recognised entity	Float		

## Post-Level Variables

Variable Name	Description	Type	Software Dependency (Python Module)	Software Dependency (Module Function)
<i>ID</i>	A unique pseudonymised ID for the respondent who authored the post	String		
<i>post_type</i>	Is it an original authored tweet, retweet, quote tweet or a reply?	String		

**Author Tweet:** This is metadata about the *author's* tweet (respondent).

<i>created_date</i>	Date the author's tweet was posted	String	datetime	date().strftime()
<i>post_source</i>	Digital source of author's posting	String		
<i>post_language</i>	Language of the author's post	String		
<i>retweet_count</i>	Total number of times the author's post was retweeted	Integer		
<i>reply_count</i>	Total number of times the author's post was replied to	Integer		
<i>like_count</i>	Total number of times the author's post was liked/favourited	Integer		
<i>quote_count</i>	Total number of times the author's post was quote tweeted	Integer		
<i>total_engagement_count</i>	Total amount of engagement the author's post received	Integer		
<i>reply_ratio</i>	Total number of replies divided by the total number of retweets and likes	Integer		
<i>contains_media</i>	Does the post contain a media attachment (Image/Video/GIF)?	Boolean		
<i>contains_geotag</i>	Does the post contain a geotag?	Boolean		
<i>reply_settings</i>	What reply settings did the author set for the post?	String		
<b>Tweet Text:</b> These variables are derived from the post text field				

<i>TEXT_post_length</i>	Total number of characters	Integer	built-in	<code>len()</code>
The following variables were applied to pre-processed text. This involves the removal of URLs, emojis, hashtags, @mentions and extra whitespace. This is conducted using a custom-built <code>strip_text()</code> function which relies on the <code>re</code> and <code>emoji</code> modules. (Excl. the hashtag, mention and emoji count variables)				
<i>TEXT_character_length</i>	Total number of textual characters	Integer	<code>textstat</code>	<code>char_count()</code>
<i>TEXT_capt_character_count</i>	Total number of capitalised characters	Integer	built-in	<code>is.upper()</code>
<i>TEXT_capt_character_prop</i>	Proportion of capitalised characters to total number of characters	Float		
<i>TEXT_letter_count</i>	Total number of letters	Integer	<code>textstat</code>	<code>letter_count()</code>
<i>TEXT_word_count</i>	Total number of words	Integer	<code>textstat</code>	<code>lexicon_count()</code>
<i>TEXT_unique_word_count</i>	Total number of unique words	Integer	built-in	<code>set()</code>
<i>TEXT_sentence_count</i>	Total number of sentences	Integer	<code>textstat</code>	<code>sentence_count()</code>
<i>TEXT_syllable_count</i>	Total number of syllables	Integer	<code>textstat</code>	<code>syllable_count()</code>
<i>TEXT_monosyllable_count</i>	Total number of monosyllables	Integer	<code>textstat</code>	<code>monosyllabcount()</code>
<i>TEXT_polysyllable_count</i>	Total number of polysyllables	Integer	<code>textstat</code>	<code>polysyllabcount()</code>
<i>TEXT_grammar_score</i>	Proportion of grammatical errors to total number of words	Float	<code>language_tool_python</code>	<code>tool_check()</code>
<i>TEXT_reading_ease</i>	Flesch Reading Ease Score. The higher the score, the easier the readability of the text. (Max 121.22)	Float	<code>textstat</code>	<code>flesch_reading_ease()</code>
<i>TEXT_non_eng_reading_ease</i>	McAlpine EFLAW Readability Score. Readability of an English text for a foreign learner of English. It is recommended to aim for a score equal to or lower than 25	Float	<code>textstat</code>	<code>mcalpine_eflaw()</code>
<i>TEXT_reading_grade</i>	Estimated school grade level required to understand the text. Consensus score based on seven separate readability tests	Float	<code>textstat</code>	<code>text_standard()</code>
<i>TEXT_exclamation_count</i>	Total number of exclamation marks	Integer	built-in	<code>count()</code>
<i>TEXT_question_mark_count</i>	Total number of question marks	Integer	built-in	<code>count()</code>
<i>TEXT_hashtag_count</i>	Total number of hashtags	Integer	<code>re</code>	<code>findall()</code>
<i>TEXT_emoji_count</i>	Total number of emojis	Integer	<code>emoji</code>	<code>emoji_count()</code>

<code>TEXT_mentions_count</code>	Total number of @mentions	Integer	<code>re</code>	<code>findall()</code>
<code>TEXT_url_count</code>	Total number of URLs	Integer	<code>re.findall()</code>	

Twitter/X reformats in-text URLs to its own domain: <https://t.co/>. To find the original URL domain, the full URL listed in the `extended_urls` column must be used. The domain of URL is extracted using the `extract().domain` function from the `tldextract` module and matched to a list of domains in a pre-built CSV file scrubbed from the Media Bias/Fact Check website: <https://mediabiasfactcheck.com/>. This contains 1579 unique domains for major media sources around the world, as well as scores of their ideological bias and reporting credibility.

<code>TEXT_url_mbfc_matches</code>	Total number of URL domain matches against the Media Bias/Fact Check (MB/FC) database	Integer	built-in	<code>Isin()</code>
<code>TEXT_url_mbfc_bias</code>	Media Bias/Fact Check bias score(s) for URL domain matches	Integer		
<code>TEXT_url_mbfc_credibility</code>	Media Bias/Fact Check factual reporting credibility score(s) for URL domain matches	Integer		

The following variables are estimated using pre-trained models from the *TweetNLP* module: <https://github.com/cardiffnlp/tweetnlp>

<code>TEXT_topics</code>	Topic classification label(s)	String	<code>tweetnlp.load_model('topic_classification')</code>	<code>topic()</code>
<code>TEXT_topic_prob</code>	Certainty estimates for each topic label	Float		
<code>TEXT_sentiment</code>	Sentiment classification label	String	<code>tweetnlp.load_model('sentiment')</code>	<code>sentiment()</code>
<code>TEXT_sentiment_prob</code>	Certainty estimate for the sentiment classification	Float		
<code>TEXT_irony</code>	Irony classification label	String	<code>tweetnlp.load_model('irony')</code>	<code>irony()</code>
<code>TEXT_irony_prob</code>	Certainty estimate for the irony classification	Float		
<code>TEXT_offensive</code>	Offensive language classification label	String	<code>tweetnlp.load_model('offensive')</code>	<code>offensive()</code>
<code>TEXT_offensive_prob</code>	Certainty estimate for the offensive language classification	Float		

<i>TEXT_emotion</i>	Emotion classification label	String	tweetnlp.load_model('emotion')	emotion()
<i>TEXT_emotion_prob</i>	Certainty estimate for the emotion classification	Float		
<i>TEXT_hate</i>	Hate speech classification label	String	tweetnlp.load_model('hate')	hate()
<i>TEXT_hate_prob</i>	Certainty estimate for the hate speech classification	Float		
<i>TEXT_entity_count</i>	Total number of named entities recognised	Integer	tweetnlp.load_model('ner')	ner()

**Referenced Tweet: (RT)** – If the tweet is a retweet, quote tweet, or a reply, the API also returns information about the original tweet. This is metadata about the *original* (referenced) tweet.

<i>RT_referenced_author_ID</i>	Pseudonymised ID of the author of the original tweet	String		
<i>RT_created_date</i>	Date the original tweet was posted	String		
<i>RT_post_source</i>	Digital source of original posting	String		
<i>RT_post_language</i>	Language of the original post	String		
<i>RT_retweet_count</i>	Total number of times the original post was retweeted	Integer		
<i>RT_reply_count</i>	Total number of times the original post was replied to	Integer		
<i>RT_like_count</i>	Total number of times the original post was liked/favourited	Integer		
<i>RT_quote_count</i>	Total number of times the original post was quote tweeted	Integer		
<i>RT_total_engagement_count</i>	Total amount of engagement the original post received	Integer		
<i>RT_reply_ratio</i>	Total number of replies divided by the total number of retweets and likes	Integer		
<i>RT_contains_media</i>	Does the original post contain a media attachment (Image/Video/GIF)?	Boolean		
<i>RT_contains_geotag</i>	Does the original post contain a geotag?	Boolean		